

Understanding Hadoop, HDFS, and MapReduce

Let's break down the key concepts and mechanisms of **Hadoop**, **HDFS**, **MapReduce**, and the related components in detail with examples to provide a clear understanding.

1. Hadoop and Its Versions

Hadoop is an open-source framework for processing and storing vast amounts of data in a distributed environment. It is designed to scale up from a single server to thousands of machines, offering high fault tolerance.

Hadoop Versions:

- **Hadoop 1.x:** First major version that introduced the basic MapReduce framework, HDFS, and YARN as a resource manager.
- **Hadoop 2.x:** Introduced **YARN (Yet Another Resource Negotiator)** for resource management, enabling better scalability, flexibility, and support for non-MapReduce applications.
- **Hadoop 3.x:** Improved scalability, performance, and added new features like support for GPUs, erasure coding, and improvements in HDFS, YARN, and other modules.

2. Hadoop Framework Structure

HDFS (Hadoop Distributed File System)

HDFS is the distributed file system that allows Hadoop to store data across many machines. It is designed for high throughput and fault tolerance.

- **Block Size:** By default, HDFS splits large files into fixed-size blocks (128MB or 256MB), and these blocks are distributed across different nodes.
 - **Replication:** Blocks are replicated for fault tolerance (default replication factor: 3).
-

MapReduce

MapReduce is a programming model used for processing and generating large datasets. It consists of two main phases:

- **Map Phase:** The input data is divided into small chunks, and a map function processes each chunk in parallel.
- **Reduce Phase:** After the map phase, the output from all the mappers is grouped, and a reduce function consolidates this data.

3. Key Hadoop Components

NameNode (Master Node)

- The **NameNode** is the central server in HDFS that manages the **metadata**. It stores information about which blocks are stored on which nodes and manages the file system namespace. However, it does **not store data**.
- **Primary Function:** Ensures that file system operations like opening, closing, and renaming files are properly managed.

DataNode (Slave Node)

- **DataNodes** are the worker nodes in HDFS. They store the actual **data** in the form of blocks on their local disks.
 - **Primary Function:** Store, retrieve, and periodically send heartbeats to the NameNode to report their health and status.
-

JobTracker (Resource Manager in YARN)

- In **Hadoop 1.x**, the **JobTracker** was responsible for managing jobs and allocating resources for MapReduce jobs.
- In **Hadoop 2.x and later**, **YARN** (Yet Another Resource Negotiator) replaces the JobTracker and acts as the resource manager to handle resource allocation across multiple applications.

TaskTracker (NodeManager in YARN)

- In **Hadoop 1.x**, **TaskTrackers** were responsible for executing individual tasks (map and reduce tasks).
 - In **Hadoop 2.x and beyond**, **NodeManagers** take over the role of TaskTrackers, managing resources and ensuring the execution of tasks in the cluster.
-

4. HDFS Workflow

1. File Write:

- When a client writes a file, the file is split into blocks.
- These blocks are stored across various **DataNodes** based on the replication factor set in HDFS.
- The **NameNode** keeps track of which DataNodes store each block and its replicas.

2. File Read:

- When a client reads a file, it contacts the **NameNode** for metadata to find where the blocks are stored.
 - It then directly contacts the **DataNodes** to fetch the blocks.
-

5. Replication Factor and Fault Tolerance

Replication Factor

- **Definition:** The replication factor is the number of copies of each data block that HDFS maintains across different DataNodes to ensure data availability and fault tolerance.
- **Default Value:** 3 (meaning 3 copies of each block are stored).

Why Replication?

- **Fault Tolerance:** If one node fails, there are multiple replicas available for data recovery.
 - **Data Availability:** Multiple replicas allow for quicker data retrieval as requests can be served by any DataNode holding the replica.
-

Example:

1. **Cluster Size (N):** 10 nodes
2. **Replication Factor (R):** 3

Each block of data will be replicated on 3 different nodes. If one node fails, the data can still be accessed from the other 2 nodes.

6. Ideal Replication Factor

How to Choose Ideal Replication Factor:

- The replication factor depends on the **cluster size**, **fault tolerance** requirements, and the **criticality** of the data.
- The formula is often used for choosing the replication factor based on fault tolerance:

$$R = N \times F$$

Where:

- R : Replication factor
 - N : Number of nodes in the cluster
 - F : Fault tolerance factor (the fraction of nodes that can fail)
-

Example:

Cluster Size: 100 nodes

Fault Tolerance: 10% (i.e., 10 nodes can fail)

Ideal Replication Factor:

$$R = 100 \times 0.1 = 10$$

Thus, a replication factor of 10 would ensure data is always accessible, even if 10 nodes fail.

7. Heart Signal in HDFS

The **heartbeat signal** is a mechanism used by **DataNodes** to communicate with the **NameNode**.

- **What it consists of:**
 - A **heartbeat** is a regular signal sent from a **DataNode** to the **NameNode** to indicate its health and status.
 - It contains information such as:
 - Disk usage
 - Block storage details
 - Node health status
 - **Characteristics:**
 - Sent every few seconds (default: 3 seconds).
 - It's a lightweight message to ensure the **NameNode** knows which **DataNodes** are alive.
 - **Importance:**
 - If the **NameNode** doesn't receive a heartbeat from a **DataNode** for a certain period (e.g., 10 minutes), it assumes the **DataNode** has failed.
 - This triggers replication of blocks to maintain the desired replication factor.
-

8. Negative Scenarios in Hadoop (Failures and Recovery)

Under-Replication:

- **Scenario:** A DataNode fails, and the block count drops below the replication factor.
- **Example:** A block has 3 replicas (Replication Factor = 3), but after a node failure, there are only 2 replicas left.
- **Action:** The NameNode detects under-replication and directs other DataNodes to create another replica of the missing block.

Over-Replication:

- **Scenario:** A DataNode recovers and starts storing a block, causing the replication to exceed the desired number of replicas.
- **Action:** The NameNode detects over-replication and deletes excess replicas, ensuring the replication factor is maintained.

Algorithm for Shortest Path:

- When data needs to be replicated or retrieved, the **NameNode** may choose the **shortest path** or the **most quickly reachable node** based on the network topology, ensuring efficient data access.
-

Summary with Example:

Imagine you are working with a cluster of 10 nodes. Your data is split into blocks and stored across these nodes. You set a **replication factor** of 3, meaning each block has 3 copies stored on different nodes. If one node fails (e.g., Node 2), the system still has 2 copies of the block, ensuring data availability. The **heartbeat signal** keeps the system aware of which nodes are active, and if Node 2 goes offline, the system triggers replication to a new node to restore the replication factor.

This structure ensures that Hadoop provides **fault tolerance**, **scalability**, and **reliability** while processing vast amounts of data across distributed systems.